

# Cat or Dog?

Aaron Rubenacker, Rita  
Tong, Tai Yu Pan

# Project overview

- Binary image classification :
  - cat
  - dog
- Data available on Kaggle website
- Two models are implemented in this project



<http://spcasuncoast.org/wp-content/uploads/2017/02/?SD>

# Softmax

- Logistic regression function
- The function contains K-dimension vector, and all entries sum up to 1

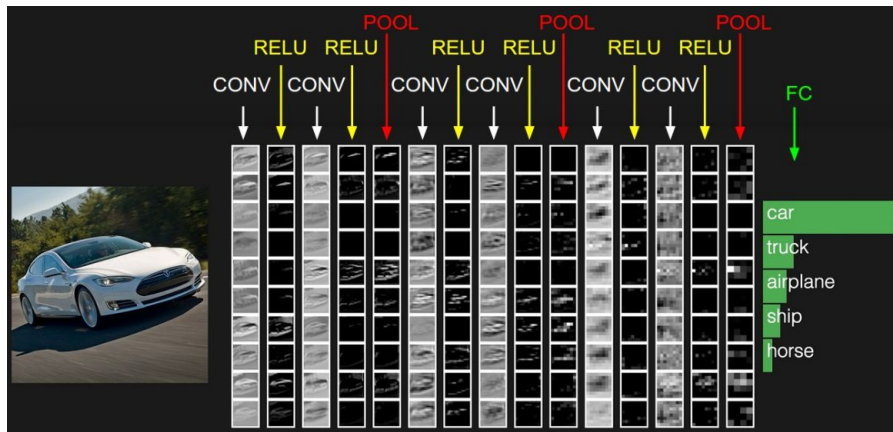
$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K.$$

- Multi-class classification

# Convolutional Neural Networks

There are four basic layers in CNN:

1. Convolution layer
2. ReLU function
3. Pooling layer
4. Fully connected layer



<http://cs231n.github.io/convolutional-networks/>

# First Model – Softmax

## Code

```
# Define the classifier's result
logits = tf.matmul(images_placeholder, weights) + biases

# Define the loss function
loss =
tf.reduce_mean(tf.nn.sparse_softmax_cross_entropy_with_logits(logits=log
its, labels=labels_placeholder))

# Define the training operation
train_step =
tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)

# Operation comparing prediction with true label
correct_prediction = tf.equal(tf.argmax(logits, 1), labels_placeholder)

# Operation calculating the accuracy of our predictions
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
```

## Result

```
Step    0: training accuracy 0.5
Step   100: training accuracy 0.49
Step   200: training accuracy 0.6
Step   300: training accuracy 0.45
Step   400: training accuracy 0.47
Step   500: training accuracy 0.5
Step   600: training accuracy 0.47
Step   700: training accuracy 0.57
Step   800: training accuracy 0.45
Step   900: training accuracy 0.51
Test accuracy 0.5554
Total time: 36.22s
```


# Reasons for a low accuracy

1. Colors
2. Features
3. The place of pixels

# Second Model - CNN

## Case Study: LeNet-5

[LeCun et al., 1998]




Conv filters were 5x5, applied at stride 1  
Subsampling (Pooling) layers were 2x2 applied at stride 2  
i.e. architecture is [CONV-POOL-CONV-POOL-CONV-FC]

<http://cs231n.github.io/convolutional-networks/>

# Second Model - CNN

## Case Study: AlexNet

[Krizhevsky et al. 2012]



Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1

[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

[4096] **FC7**: 4096 neurons

[1000] **FC8**: 1000 neurons (class scores)

### Details/Retrospectives:

- first use of ReLU
- used Norm layers (not common anymore)
- heavy data augmentation
- dropout 0.5
- batch size 128
- SGD Momentum 0.9
- Learning rate 1e-2, reduced by 10 manually when val accuracy plateaus
- L2 weight decay 5e-4
- 7 CNN ensemble: 18.2% -> 15.4%

<http://cs231n.github.io/convolutional-networks/>



# Second Model - CNN

## Summary

- ConvNets stack CONV, POOL, FC layers
- Trend towards smaller filters and deeper architectures
- Trend towards getting rid of POOL/FC layers (just CONV)
- Typical architectures look like

**$[(\text{CONV-RELU}) * N - \text{POOL?}] * M - (\text{FC-RELU}) * K, \text{SOFTMAX}$**

where N is usually up to ~5, M is large,  $0 \leq K \leq 2$ .

- but recent advances such as ResNet/GoogLeNet challenge this paradigm

# Second Model - CNN

- Environments:

Python == 3.5.6

tensorflow == 1.10.0

tflearn == 0.3.2

- Input image:  $32 \times 32 \times 3$

# Second Model – CNN

```
def create_model():
    convnet = input_data(shape=[None, IMG_SIZE, IMG_SIZE, 3],
                           name='input')

    conv1 = conv_2d(convnet, 32, 2, activation='relu', name='conv1')
    conv1_ = max_pool_2d(conv1, 2)

    conv2 = conv_2d(conv1_, 64, 2, activation='relu', name='conv2')
    conv2_ = max_pool_2d(conv2, 2)

    conv3 = conv_2d(conv2_, 32, 2, activation='relu', name='conv3')
    conv3_ = max_pool_2d(conv3, 2)

    conv4 = conv_2d(conv3_, 64, 2, activation='relu', name='conv4')
    conv4_ = max_pool_2d(conv4, 2)

    conv5 = conv_2d(conv4_, 32, 2, activation='relu', name='conv5')
    conv5_ = max_pool_2d(conv5, 2)

    conv6 = conv_2d(conv5_, 64, 2, activation='relu', name='conv6')
    conv6_ = max_pool_2d(conv6, 2)

    fc1 = fully_connected(conv6_, 1024, activation='relu', name='fc1')
    fc1_ = dropout(fc1, 0.8)

    fc2 = fully_connected(fc1_, 2, activation='softmax', name='fc2')
    fc2_ = regression(fc2, optimizer='adam', learning_rate=LR,
                      loss='categorical_crossentropy', name='targets')

    model = tflearn.DNN(fc2_, tensorboard_dir='log')
    return model
```

# Second Model - CNN


```
Training Step: 61500 | total loss: 0.06957 | time: 88.467s
| Adam | epoch: 196 | loss: 0.06957 - acc: 0.9804 | val_loss: 0.05519 - val_acc: 0.9812 -- iter: 09728/20000
--
Training Step: 61661 | total loss: 0.03424 | time: 179.967s
| Adam | epoch: 196 | loss: 0.03424 - acc: 0.9900 | val_loss: 0.04190 - val_acc: 0.9888 -- iter: 20000/20000
--
Training Step: 61974 | total loss: 0.05505 | time: 145.492s
| Adam | epoch: 197 | loss: 0.05505 - acc: 0.9888 | val_loss: 0.02407 - val_acc: 0.9910 -- iter: 20000/20000
--
Training Step: 62000 | total loss: 0.02000 | time: 43.913s
| Adam | epoch: 198 | loss: 0.02000 - acc: 0.9942 | val_loss: 0.02151 - val_acc: 0.9922 -- iter: 01664/20000
--
Training Step: 62287 | total loss: 0.03045 | time: 180.257s
| Adam | epoch: 198 | loss: 0.03045 - acc: 0.9914 | val_loss: 0.02636 - val_acc: 0.9919 -- iter: 20000/20000
--
Training Step: 62500 | total loss: 0.03250 | time: 109.991s
| Adam | epoch: 199 | loss: 0.03250 - acc: 0.9860 | val_loss: 0.03269 - val_acc: 0.9849 -- iter: 13632/20000
--
Training Step: 62600 | total loss: 0.01942 | time: 180.209s
| Adam | epoch: 199 | loss: 0.01942 - acc: 0.9927 | val_loss: 0.02158 - val_acc: 0.9936 -- iter: 20000/20000
--
Training Step: 62913 | total loss: 0.03722 | time: 145.902s
| Adam | epoch: 200 | loss: 0.03722 - acc: 0.9833 | val_loss: 0.02176 - val_acc: 0.9935 -- iter: 20000/20000
```

Computer Environment:  
model == MacBook Pro (15-inch,  
2016)  
processor == 2.9 GHz Intel Core i7  
Graphics == Radeon Pro 460 4096  
MB  
It takes in total 16 hrs 33 mins to  
train the model for 200 epoch  
containing 62913 steps.


# Second Model

Accuracy vs steps on the training set and test set respectively:

Accuracy




Accuracy/Validation




Loss values vs steps on training set and test set respectively:

Loss

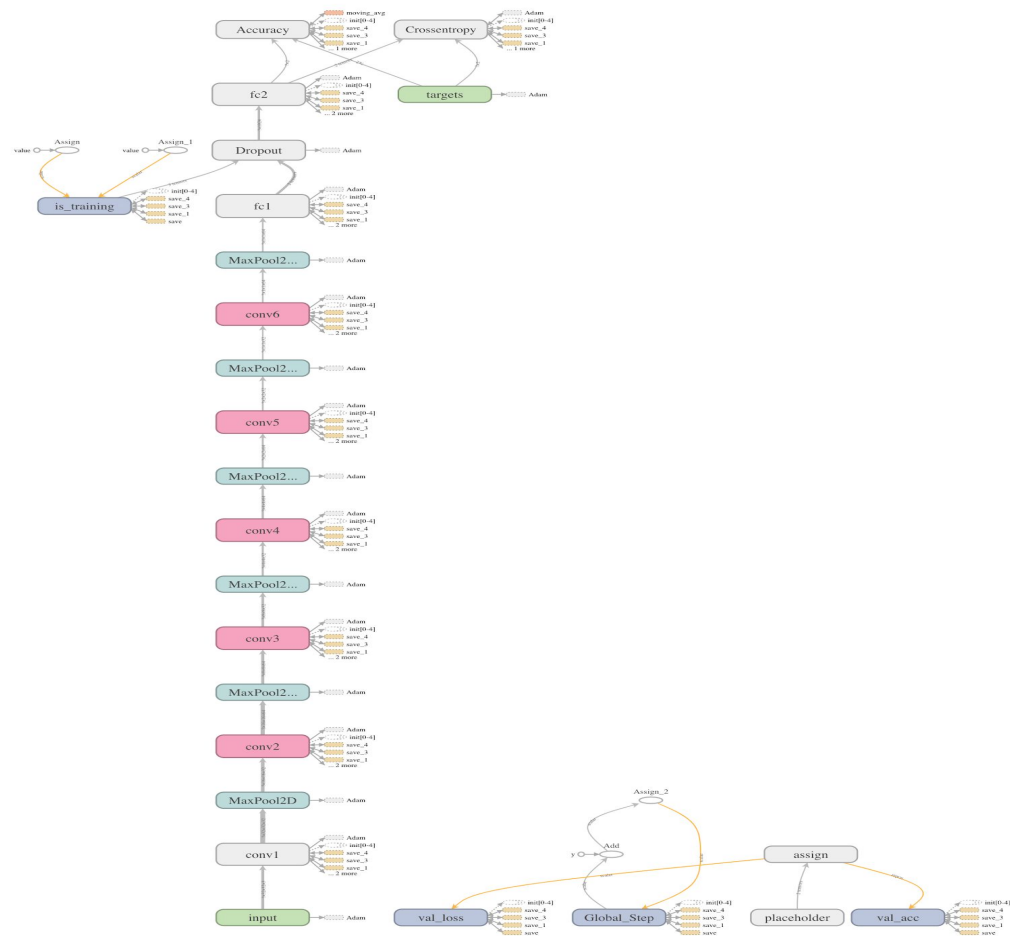


Loss/Validation



# Evaluate

## CNN Model



# Result from CNN

1. Solve problems from Softmax
  - colors
  - features
  - place of pixels
2. 99% accuracy
3. Small images



# Evaluation

Sample output:

ID: 1010, Dog



ID: 11854, Cat



ID: 3993, Cat



ID: 8271, Dog



ID: 2268, Cat



ID: 9710, Cat



ID: 8661, Cat



ID: 4298, Dog



ID: 6164, Cat



ID: 2480, Cat



ID: 5910, Cat



ID: 10003, Dog





Demo

Thank you for listening